

Comparing Performance of KMP-Algorithm in Serial and Parallel Environments

Joel John Thomas
USN: 1MS18CS053

Abstract

In order to make the most out of the CPU's resources, we must find different ways to implement the algorithms. KMP (the Knuth–Morris–Pratt) algorithm is used **to find a "Pattern" in a "Text"**. This algorithm compares character by character from left to right. But whenever a mismatch occurs, it uses a preprocessed table called "Prefix Table" to skip characters comparison while matching. To activate parallel programming, we will use MPI clustering programming. MPI is a message-passing model of distributed memory. To assess parallel programming's performance, we compare it to a sequential KMP algorithm, which does not make efficient use of the CPU's resources. To determine whether an algorithm is more efficient, we evaluate the performance of sequential and parallel KMP algorithms.

Introduction

The exponential growth of processing and network speeds means that parallel architecture isn't just a good idea; it's necessary. Big data and the IoT will soon force us to crunch trillions of data points at once. Nowadays new applications always need faster processors in today's world to operate efficiently. Parallel systems are employed in a lot of business applications nowadays. Parallel computing uses multiple computer cores to attack several operations at once. Unlike serial computing, parallel architecture can break down a job into its component parts and multi-task them. Parallel computer systems are well suited to modeling and simulating real-world phenomena. Without parallel computing, performing digital tasks would be tedious, to say the least. The advantages of parallel computing are that computers can execute code more efficiently, which can save time and money by sorting through "big data" faster than ever.

Parallel systems, divide the programs into multiple fragments and process them all at the same time in order to reduce the execution time of the program. By

constructing a parallel processing bundle or a combination of both entities, a parallel system can deal with many processors, machines, computers, or CPUs, among other things.

String searching is one of the most common problems where it tries to find whether a string is a sub-string of another particular string or not. There are several algorithms to solve this problem. These algorithms are mostly used in text edit processing, image processing, literature retrieval, natural language recognition, spell checking, and many other fields.

The most common string searching algorithm is Knuth–Morris–Pratt (KMP) algorithm. It is a prefix based searching algorithm faster than normal string searching algorithm like 'Naive String Searching Algorithm'. The algorithm uses degenerating property (pattern having same sub-patterns appearing more than once in the pattern) of the pattern and improves the worst case complexity to $O(n)$.

We can make this algorithm faster if we make it parallelly executable. As it is linear time complexity algorithm, it will consume more time if the text data is big. If we run the KMP algorithm on multiple computers to handle larger data, it will be less time-consuming. By doing that, it can be used in large-scale applications. To achieve parallelism in KMP algorithm, we will use MPI clustering programming. MPI is a message-passing model of distributed memory. The aim of the MPI is to establish an efficient and flexible standard for message passing which is used for writing message passing programs. As such, MPI is the vendor independent and having message passing library. The benefit of developing message passing software using MPI helps to achieve design goals like portability, efficiency and flexibility. MPI performs their operations by using the methods such as `MPI_send`, `MPI_receive`, `MPI_COMM_WORLD`, `MPI_Init` and many more by Using these kind of functionality the processes communicate with themselves.

The steps of parallel implementations using MPI are given below:

- 1) Read the input file of Snort rules, prepare the input stream of that file, obtain the file size, allocate memory to contain whole file, copy the file into the buffer and divide it across processes using MPI_Bcast and MPI_Barrier.
- 2) Prepare the input stream of pattern file and get the number of lines.
- 3) Divide the input file among the processors.
- 4) Call the KMP search function and compute the longest proper suffix and at last show the matching patterns as output.

Code:

```
1
2 #include "mpi.h"
3 #include <stdlib.h>
4 #include <stdio.h>
5 #include <string.h>
6 #include <malloc.h>
7
8 #define master 0
9 #define MAX(a,b) ((a) > (b) ? a : b)
10 #define MILLION 1000000L
11
12 #define bug printf("trust\n");
13
14 int* kmp(char* target, char* pattern, int* table, int myrank);
15 int* kmptable(char* pattern, int len);
16 void print(char* arr, int id, int len);
17 void fillup(char** matrix, int rankid, char* msg, int len, int offset);
18 void freeDouble(char** matrix, int nproc);
19 void printMatrix(char **matrix, int nproc, int id, int length);
20 int* getRealIdx(int* answer, int len, int myrank, int x, int y, int* length);
21 void pinpoint(int* result, int* msg, int shortlen);
22
23
24 void pinpoint(int* result, int* msg, int shortlen){
25     int j = 0;
26     for(j=0; j<shortlen; j++){
27         if(j > 0 && msg[j] == 0)
28             break;
29         else if(j == 0 && msg[j] == 0)
30             continue;
31         else{
32             result[msg[j]] = 1;
33         }
34     }
35 }
36
37 // get the exact idx(es) in each process
38 // return back an int pointer
39 int* getRealIdx(int* answer, int len, int myrank, int x, int y, int* length){
40     int* package = (int *) calloc(len, sizeof(int));
41     int index;
42     int real_len = 0;
43     int i;
44     for(i=0; i < len; i++){
45         if(i > 0 && answer[i] == 0)
46             break;
47         else if(i == 0 && answer[i] == 0)
48             continue;
49         else{
50             if(myrank == 0){
51                 index = answer[i];
52                 package[real_len++] = index;
53                 // printf("This is REAL index: %d\n", index);
54             }else{
```



```

156 char pattern[] = "pryyaogqzekpumxktopubehtpzeqlkfcifewmwsygnotkbmyqhzqujpayflinpsuvqisobabtrtsnjrofhkpkpgvxtmpqifycubcxvwmngayxpegbvfnkn";
157
158 n = strlen(target);
159 m = strlen(pattern);
160
161 int tag = 1;
162 int tag2 = 2;
163
164 int* kmp_table = kmptable(pattern, m);
165
166
167
168 int myrank, nproc;
169 MPI_Init(&argc, &argv); //Initialize the MPI
170 MPI_Comm_rank(MPI_COMM_WORLD, &myrank); //number of processes that given in command line
171 MPI_Comm_size(MPI_COMM_WORLD, &nproc); //ranks the process rank
172
173
174
175
176 MPI_Status status;
177
178 int start_pos = n/nproc;
179 int str_per_proc = n/nproc;
180 int end = str_per_proc + n % nproc;
181
182
183 char** matrix = (char**)malloc(nproc*sizeof(char*));
184 for(i=0; i<nproc; i++){
185     matrix[i] = (char*)malloc((end-m+1) * sizeof(char));
186 }
187 printMatrix(matrix, nproc, myrank, 10);
188
189 char send_msg[MAX(m-1, end)]; // 4-1 = 3
190 char recv_msg[MAX(m-1, end)];
191 char end_msg[MAX(m-1, end)];
192
193 double start, stop;
194
195
196 if(myrank == master){
197     printf("Result KMP Sequential\n");
198     start = MPI_Wtime();
199     for(i=1; i< nproc; i++){
200         if(i == nproc-1){
201             strncpy(send_msg, target + i*start_pos, end);
202             MPI_Send(send_msg, end, MPI_CHAR, i, tag, MPI_COMM_WORLD);
203         }else{
204             strncpy(send_msg, target + i*start_pos, str_per_proc);
205             MPI_Send(send_msg, str_per_proc, MPI_CHAR, i, tag, MPI_COMM_WORLD);
206         }
207     }
208 }
209

```

```

208
209 }
210 strncpy(send_msg, target + master * start_pos, str_per_proc);
211 fillup(matrix, myrank, send_msg, str_per_proc, 0);
212
213 // The last process
214 }else if(myrank == nproc-1){
215     MPI_Recv(recv_msg, end, MPI_CHAR, master, tag, MPI_COMM_WORLD, &status);
216
217     fillup(matrix, myrank, recv_msg, end, m-1);
218 }
219 else{
220     MPI_Recv(recv_msg, str_per_proc, MPI_CHAR, master, tag, MPI_COMM_WORLD, &status);
221     fillup(matrix, myrank, recv_msg, str_per_proc, m-1);
222 }
223
224 MPI_Barrier(MPI_COMM_WORLD);
225
226 if(myrank == nproc-1){
227
228     MPI_Recv(end_msg, m-1, MPI_CHAR, myrank-1, tag, MPI_COMM_WORLD, &status);
229
230     fillup(matrix, myrank, end_msg, m-1, 0);
231     int* answer = kmp(matrix[myrank], pattern, kmp_table, myrank);
232
233     int len;
234     // -----
235     int* result = getRealIdx(answer, end-m+1, myrank, str_per_proc, m-1, &len);
236     // -----
237     free(answer);
238
239     MPI_Send(result, len, MPI_INT, master, tag2, MPI_COMM_WORLD);
240     free(result);
241
242 }else{
243     strncpy(send_msg, target + str_per_proc-m + 1 + myrank * str_per_proc, m-1);
244     MPI_Send(send_msg, m-1, MPI_CHAR, myrank+1, tag, MPI_COMM_WORLD);
245
246     /* Processes other than master one
247     * Re-recv more msg from the previous process whose rank is myrank-1
248     * kmp implementation
249     * Send the result back to master process
250     */
251
252     if(myrank==master){
253         MPI_Recv(end_msg, m-1, MPI_CHAR, myrank-1, tag, MPI_COMM_WORLD, &status);
254         fillup(matrix, myrank, end_msg, m-1, 0);
255         // implement kmp to get the matching result
256         int* answer = kmp(matrix[myrank], pattern, kmp_table, myrank);
257         int len;
258         int* result = getRealIdx(answer, end-m+1, myrank, str_per_proc, m-1, &len);
259         free(answer);
260         MPI_Send(result, len, MPI_INT, master, tag2, MPI_COMM_WORLD);
261         free(result);

```

```

271         int* answer = kmp(matrix[myrank], pattern, kmp_table, myrank);
272         int len;
273         int* result = getRealIdx(answer, end+m-1, myrank, str_per_proc, m-1, &len);
274         pinpoint(final_result, result, len);
275         free(answer);
276         free(result);
277         int j;
278         for (j = 0; j < end+m-1; j++)
279         {
280             recv[j] = 0;
281         }
282
283         for(j = 1; j < nproc; j++){
284             MPI_Recv(recv, end+m-1, MPI_INT, j, tag2, MPI_COMM_WORLD, &status);
285             pinpoint(final_result, recv, end+m-1);
286         }
287         stop = MPI_Wtime();
288         free(recv);
289         double timeDifference = (stop-start) * MILLION;
290         printf("Target Length: %d\n", n);
291         printf("Pattern Length: %d\n", m);
292         printf("Time Elapsed: %0.05f\n", timeDifference);
293         // for(j=0; j<n; j++){
294         //     if(final_result[j] == 1)
295         //         printf("Find a matching substring starting at: %d.\n", j);
296         // }
297         printf("\n");
298         free(final_result);
299     }
300 }
301
302 MPI_Barrier(MPI_COMM_WORLD);
303 freeDouble(matrix, nproc);
304 MPI_Finalize();
305
306 return 0;
307 }
308
309 int* kmpTable(char* pattern, int len){
310     int k = 0;
311     int i = 1;
312     int* table = (int*)malloc(len * sizeof(int));
313     table[0] = k;
314     for(i=1; i<len; i++){
315         while(k > 0 && pattern[i-1] != pattern[i]){
316             table[i] = k-1;
317             k = table[i];
318         }
319         if(pattern[i] == pattern[k])
320             k++;
321         table[i] = k;
322     }
323     return table;
324 }

```

Output:

On a Computer without using MPI

Result KMP_MPI

Start Point: 49521.

Target Length: 99044

Pattern Length: 122

Time Elapsed: 2581.42600

On both Computer using MPI Clustering

Result KMP_MPI

Start Point: 49521.

Target Length: 99044

Pattern Length: 122

Time Elapsed: 698.22000

```

data1000000.txt kmp_MPI.c pattern.txt test.txt
data10000.txt kmp_MPI.out Temporary.c
joel@joel-Inspiron-3576:~/Downloads/KMP-Parallel-Using-MPI-master$ mpicc -g -Wall -o test kmp_MPI.c
kmp_MPI.c: In function 'main':
kmp_MPI.c:153:9: warning: unused variable 'diff' [-Wunused-variable]
  153 |     double diff;
      |           ^~~~~
kmp_MPI.c:152:27: warning: unused variable 'endi' [-Wunused-variable]
  152 |     struct timespec start1, end1;
      |                               ^~~~~
kmp_MPI.c:152:19: warning: unused variable 'start1' [-Wunused-variable]
  152 |     struct timespec start1, end1;
      |                   ^~~~~~
kmp_MPI.c:148:8: warning: unused variable 'j' [-Wunused-variable]
  148 |     int i, j;
      |            ^
joel@joel-Inspiron-3576:~/Downloads/KMP-Parallel-Using-MPI-master$ mpiexec -n 2 ./test
Invalid MIT-MAGIC-COOKIE-1 key

Result KMP Sequential
Start Point: 49521.
Target Length: 99044
Pattern Length: 122
Time Elapsed: 2581.42600

joel@joel-Inspiron-3576:~/Downloads/KMP-Parallel-Using-MPI-master$ mpicc -O test kmp_MPI.c
gcc: error: unrecognized command line option '-O'
joel@joel-Inspiron-3576:~/Downloads/KMP-Parallel-Using-MPI-master$ mpicc -o test kmp_MPI.c
joel@joel-Inspiron-3576:~/Downloads/KMP-Parallel-Using-MPI-master$ mpirun -np 2 ./test
Invalid MIT-MAGIC-COOKIE-1 key

Result KMP Sequential

Start Point: 49521.
Target Length: 99044
Pattern Length: 122
Time Elapsed: 698.22000

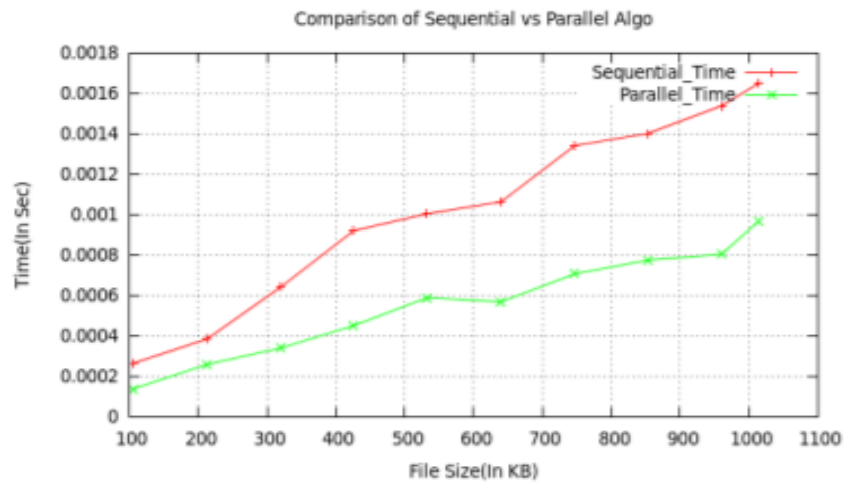
joel@joel-Inspiron-3576:~/Downloads/KMP-Parallel-Using-MPI-master$ ls
data1000000.txt  HowToRun.txt      KMPSequential  test
data1000000.txt  IntroductionReport.pdf KMPSequential.c testmake.py
data100000.txt   kmp_MPI.c         pattern.txt     text.txt
data10000.txt    kmp_MPI.out       Temporary.c
joel@joel-Inspiron-3576:~/Downloads/KMP-Parallel-Using-MPI-master$ mpicc -g -Wall -o test KMPSequential.c
KMPSequential.c: In function 'main':
KMPSequential.c:17:8: warning: unused variable 'j' [-Wunused-variable]
   17 |     int i, j;
      |            ^
KMPSequential.c:17:6: warning: unused variable 'i' [-Wunused-variable]
   17 |     int i, j;
      |      ^
KMPSequential.c: In function 'kmptable':
KMPSequential.c:17:8: warning: unused variable 'j' [-Wunused-variable]
   17 |     int i, j;
      |            ^

```

Result

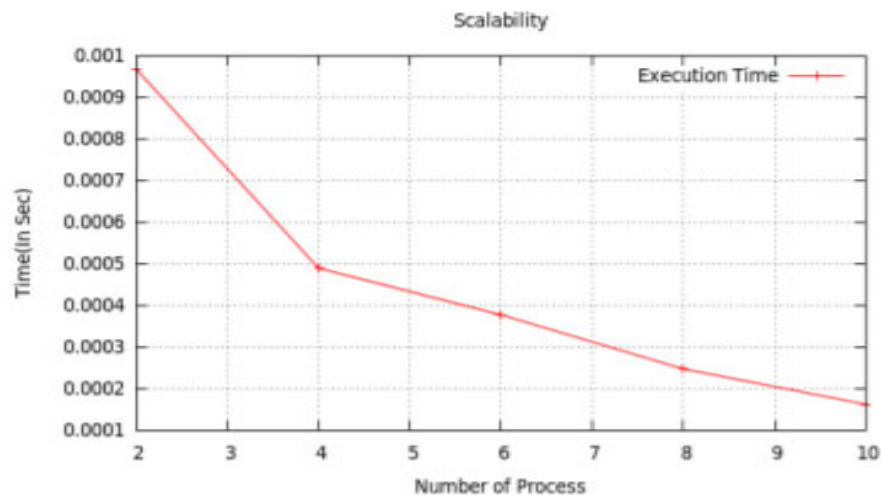
The most time consuming part of KMP search is the LPS that is Longest Prefix Suffix. The following results show the Compute time of this algorithm.

Sequential Vs Parallel Time Comparison graph:



The above graph shows us the comparison of sequential and parallel time in which we have taken the input file in KB on X-axis and Time(s) to compute in Y-axis by keeping the pattern file constant. So the above graph shows that the parallel computation time is more efficient than the sequential time and the performance of parallel is completely outplayed the sequential performance.

Scalability graph:



The above shows us the number of processes on X-axis and the time taken by that process on Y-axis. In the above graph 10 processes are taken and computed the time taken by them in seconds. It shows that if we go on increasing the number of processes, then the performance gets better. In this way the scalability is achieved.

By viewing the outputs and the graphs we can conclude by saying we implemented the Knuth-Morris-Pratt algorithm to match the Snort pattern and the performance of the parallel algorithm is far better than the sequential algorithm for multiple patterns.

References

- [1] Reetinder Sidhu and Viktor K Prasanna. Fast regular expression matching using fpgas. In Field-Programmable Custom Computing Machines, 2001. FCCM'01. The 9th Annual IEEE Symposium on, pages 227-238. IEEE, 2001.
- [2] Christopher R Clark and David E Schimmel. Scalable pattern matching for high speed networks. In Field-Programmable Custom Computing Machines, 2004. FCCM 2004. 12th Annual IEEE Symposium on, pages 249-257. IEEE, 2004.
- [3] Sarang Dharmapurikar and John Lockwood. Fast and scalable pattern matching for content filtering. In Architecture for networking and communications systems, 2005. ANCS 2005. Symposium on, pages 183-192. IEEE, 2005.
- [4] Alfred V Aho and Margaret J Corasick. Efficient string matching: an aid to bibliographic search. Communications of the ACM, 18(6):333-340, 1975.
- [5] Nen-Fu Huang, Hsien-Wei Hung, ShengHung Lai, Yen-Ming Chu, and Wen-Yen Tsai. A gpu-based multiple-pattern matching algorithm for network intrusion detection systems. Advanced Information Networking and Applications-Workshops, 2008. AINAW 2008. 22nd International Conference on, pages 62-67. IEEE, 2008.